



# Achieving Peak Performance at Predictable Cost: An Optimized Hybrid Architecture for Enterprise PDF Intelligence in India

## Comparative Analysis of Large Language Models for High-Fidelity Document Intelligence

The selection of a foundational large language model (LLM) is the most critical decision in architecting a high-fidelity document intelligence platform. This analysis evaluates both proprietary "Big-Tech" and leading open-source models against the user's primary requirements, which prioritize accuracy, legal and compliance analysis quality, and resistance to hallucination. The competitive landscape has matured significantly, with open-source models demonstrating capabilities that rival or even exceed those of their closed-source counterparts in specialized domains [75](#) [100](#). The evaluation is structured around the performance of these models in relevant benchmarks and their suitability for the specified use cases of legal and compliance analysis.

The top-tier proprietary models—GPT-5 variants, Claude series (Sonnet, Opus), and Gemini versions—continue to set industry standards for general-purpose reasoning and factual knowledge [61](#) [98](#). GPT-5, particularly its reasoning-optimized variants, excels in analytical tasks and complex problem-solving benchmarks like SWE-Bench Pro [43](#) [81](#). In one study, GPT-5 (Thinking) and GPT-5 (Auto) ranked first and second respectively in hallucination detection studies, indicating strong internal consistency and factuality [110](#). The Claude family, developed by Anthropic, demonstrates particular strength in nuanced evaluations such as contract issue spotting and playbook assessments, where Claude Sonnet 4.5 was identified as the strongest performer [3](#). Similarly, the release of Gemini 2.5 Pro in March 2025 marked an achievement of then-industry-best performance across numerous academic benchmarks, showcasing its broad capabilities [34](#). These models offer a high baseline of quality due to their massive training datasets and extensive fine-tuning, making them reliable choices for general-purpose question-answering. However, their generic nature may necessitate more sophisticated Retrieval-Augmented Generation

(RAG) pipelines to achieve peak performance in highly specialized domains like law and compliance.

In contrast, the open-source ecosystem has produced models that are not only competitive but also offer superior flexibility and cost-efficiency for specific applications. Models such as Qwen 3, DeepSeek V3/R1, Llama 4 Maverick, and Mistral Large are frequently cited as leaders in this space [42](#) [98](#). The emergence of models specifically fine-tuned for legal reasoning represents a significant trend. DescrybeLM, for instance, was reported to outperform leading general-purpose models like ChatGPT, Claude, and Gemini on bar exam questions, answering all 200 correctly where others missed between 13 and 23 [56](#) [71](#). This specialization is critical for achieving high scores in the "Legal Analysis Quality" criterion, which carries a 15% weight in the evaluation framework. Benchmarks like LawBench, which evaluates 51 different LLMs across various legal domains, and SafeLawBench, designed to test the safe alignment of models on legal tasks, provide objective measures of this capability [44](#) [72](#). The ability to self-host these open-source models allows for further customization through fine-tuning on proprietary legal documents, creating a system uniquely suited to the enterprise's needs.

A crucial aspect of model evaluation is assessing their propensity for hallucination, defined as generating content not present in the source documents it retrieved [10](#). While RAG is designed to mitigate this, its effectiveness is heavily dependent on the retrieval process [5](#). Advanced benchmarks like HalluHard, which focuses on challenging multi-turn hallucinations in high-stakes domains including legal cases, provide a rigorous test of a model's reliability [55](#) [76](#). Even state-of-the-art models can struggle, and the choice of model can have a direct impact on the final output's veracity. Research into methods like using a smaller, more reflective model (e.g., Llama 4 Maverick) to improve the accuracy of a larger model's response suggests that cross-model techniques can enhance factuality [112](#). Therefore, a model's performance on hallucination-specific benchmarks is a key indicator of its suitability for enterprise-grade compliance and legal analysis. The table below provides a comparative overview of leading models based on available information, highlighting their potential strengths for the specified use cases.

| Model Family | Key Variants                | Strengths for Document Intelligence                                                                                                                         | Relevant Benchmarks / Studies                                                                             |
|--------------|-----------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| OpenAI       | GPT-5 (Reasoning, Thinking) | Strong analytical reasoning, high factual accuracy, excellent performance on math and programming benchmarks <a href="#">43</a> <a href="#">81</a> .        | SWE-Bench Pro <a href="#">43</a> , ARC-AGI-2 reasoning <a href="#">43</a> , HalluLens <a href="#">111</a> |
| Anthropic    | Claude 4 (Opus, Sonnet)     | Strong performance in nuanced evaluations like contract analysis and playbook creation <a href="#">3</a> . Good balance of capability and cost.             | SafeLawBench <a href="#">44</a> , HalluHard <a href="#">55</a>                                            |
| Google       | Gemini 2.5 (Pro, Flash)     | Industry-leading performance on academic benchmarks at time of release <a href="#">34</a> . Strong multimodal capabilities.                                 | General Academic Benchmarks <a href="#">105</a>                                                           |
| Alibaba      | Qwen 3 (32B, 72B)           | Highly capable open-source models with strong performance in Chinese agricultural benchmark <a href="#">45</a> . Potential for fine-tuning.                 | LawBench <a href="#">72</a> , OmniCompliance-100K <a href="#">94</a>                                      |
| DeepSeek AI  | DeepSeek V3, R1             | Competitive open-source models noted for strong performance in Chinese Hallucination QA <a href="#">9</a> .                                                 | HalluQA <a href="#">9</a> , LEXam <a href="#">46</a>                                                      |
| Meta         | Llama 4 (Maverick, Scout)   | Frontier open-source models offering a powerful combination of size and efficiency. Maverick shows promise for factuality improvement <a href="#">112</a> . | HalluLens <a href="#">111</a>                                                                             |
| Mistral AI   | Mistral Large 3             | Efficient reasoning model with a focus on performance per token. Good for cost-sensitive deployments.                                                       | General Reasoning Benchmarks <a href="#">81</a>                                                           |
| Descrybe     | DescrybeLM                  | Specialized legal model that outperformed GPT, Claude, and Gemini on bar exam questions <a href="#">56</a> <a href="#">71</a> .                             | Bar Exam Questions <a href="#">71</a>                                                                     |

Ultimately, the choice depends on a trade-off between the proven, albeit generic, excellence of Big-Tech APIs and the specialized, customizable power of open-source models. For a use case weighted so heavily towards legal and compliance analysis, a fine-tuned open-source model like Qwen 3 or a specialized legal model like DescrybeLM presents a compelling argument for higher quality outputs. The ability to host such a model on-premise or in a private cloud environment directly addresses the enterprise's security and data governance concerns, providing a path to a solution that is not only high-performing but also trustworthy and compliant with emerging regulatory expectations in India.

## Architectural Blueprint for End-to-End PDF Document Processing

The success of any PDF question-answering system hinges on a robust and well-designed Retrieval-Augmented Generation (RAG) pipeline. A typical RAG pipeline consists of several stages: document ingestion, parsing and chunking, embedding, indexing in a

vector database, retrieval, re-ranking, prompt construction, and finally, generation [67](#) [86](#). Each component must be meticulously architected to ensure the final answer is accurate, relevant, and properly cited, thereby minimizing the risk of hallucination and meeting enterprise-grade standards for compliance and auditability. Given the company's need for sub-second latency and high-quality legal analysis, the architecture must prioritize precision over brute-force recall.

The ingestion workflow begins with uploading PDF documents, which are then parsed to extract text and metadata. Since OCR is not required, standard libraries can be used for this task. The next critical step is chunking, where the raw text is broken down into smaller, semantically coherent segments suitable for embedding [67](#). Simple methods like fixed-size windows often fail to respect natural document boundaries, leading to fragmented context. More advanced strategies, such as recursive character splitting or using language models to identify logical section breaks, are recommended. For legal and compliance documents, preserving paragraph-level context is paramount. Following chunking, each chunk is passed through an embedding model to convert it into a high-dimensional vector representation. The choice of embedding model is crucial, as it directly impacts the quality of semantic search during retrieval. Models like those from Cohere or open-source alternatives should be selected based on their performance on semantic similarity tasks relevant to the domain.

These vectors, along with their corresponding text chunks and metadata, are then indexed in a dedicated vector database. The selection of the vector database involves a trade-off between ease of management and cost control. Managed services like Pinecone or Azure AI Search offer simplicity and scalability but come with vendor lock-in and potentially higher costs [58](#). For a cost-efficient and secure architecture, a self-hosted solution like Qdrant or Weaviate running on the same cloud infrastructure as the LLM inference engine is often preferable. It provides greater control over data residency and operational costs. Once a query is received, the system first embeds the query and performs a similarity search against the vector database to retrieve the most relevant chunks. To further refine results, a re-ranking stage can be implemented, where a smaller, faster model re-orders the initial set of retrieved chunks based on their precise relevance to the query. This hybrid retrieval approach improves the quality of the context fed into the generative model.

The core of the RAG architecture is the generation stage, where the top-k retrieved chunks are formatted into a prompt alongside the original user query and sent to the chosen LLM. The prompt design is a science in itself, aiming to instruct the model to synthesize an answer strictly from the provided context and to cite the source chunks used. To enforce this, citation-enforced RAG architectures have been proposed, which

structure the output to explicitly reference the provenance of the information, enhancing explainability and trustworthiness [59](#) [93](#). This feature is essential for satisfying the user's requirement for "citation support" and is a key factor in building a defensible, auditable system. The entire process, from ingestion to generation, must be wrapped in a scalable and observable service layer. This includes implementing robust API gateways, caching mechanisms to reduce latency for repeated queries, and comprehensive monitoring for tracking latency, error rates, and resource utilization [23](#). Finally, a security architecture must be woven throughout this entire flow. This includes encrypting data at rest and in transit, implementing a Zero Trust access model with role-based controls, and maintaining immutable audit logs of all processing activities to ensure accountability [39](#) [64](#). Such a holistic architectural blueprint ensures that the final solution is not just fast and accurate, but also secure, compliant, and trustworthy.

## Deployment Modalities: A Comparative Total Cost of Ownership Analysis

The choice of deployment modality fundamentally shapes the economic and operational profile of the PDF question-answering platform. The analysis covers a spectrum from fully managed APIs to self-hosted cloud instances and on-premise hardware, evaluating each against the enterprise's projected workload and cost constraints. A realistic projection for 2026, considering an annual volume of approximately 36,500 queries (~100/day), is essential for determining the Total Cost of Ownership (TCO) over a three-year horizon. All cost calculations assume the India market and incorporate relevant tax considerations.

Cloud-native APIs from Big-Tech providers (e.g., OpenAI, Anthropic, Google) offer the lowest operational complexity and fastest time-to-market. The pricing is typically pay-as-you-go, charging per token for input and output, plus additional fees for embeddings and managed vector databases [13](#) [21](#). For a workload of 100 queries per day, the monthly cost would likely fall in the range of ₹5,000 to ₹20,000 INR, depending on the model's price point and the average length of queries and responses. While this is predictable on a month-to-month basis, the cumulative annual cost of approximately ₹60,000 to ₹240,000 INR makes this option less viable for long-term, large-scale use compared to self-hosted alternatives. Furthermore, sending internal company documents to third-party APIs introduces inherent security and privacy risks that must be mitigated with contractual safeguards [66](#).

Self-hosting on public clouds (AWS, Azure, GCP) or specialized GPU providers (RunPod, Lambda Labs) presents a compelling alternative for balancing cost, control, and performance. This approach involves renting virtual machines equipped with GPUs to run the LLM inference server. The primary cost driver is the GPU instance, priced hourly. For example, a single NVIDIA A100 GPU might cost around 2.50–4.00 per hour, while an H100 could be 6.00–9.00 per hour [53](#) [54](#) . To support 10-15 concurrent users with sub-10 second latency, a small cluster of mid-range GPUs (e.g., two A100s) would likely be required. Assuming 24/7 operation, the monthly infrastructure cost would be approximately 1,800–6,500 USD ( [FORMULA 3](#) 64,800 - \$234,000 USD, excluding storage, networking, and labor costs [15](#) [20](#) . This TCO becomes highly competitive when amortized over the total number of queries. Crucially, this model keeps data within the organization's controlled environment (either in a private VPC or on-premise), addressing security and compliance concerns. Specialized providers like Lambda Labs or Thunder Compute may offer up to 80% lower GPU costs than major cloud providers, which could significantly reduce the TCO if they meet performance and reliability requirements [91](#) .

On-premise deployment represents the ultimate in data sovereignty and control. The TCO here is dominated by upfront capital expenditure (CapEx) on servers and GPUs. A small enterprise deployment might require a server housing multiple RTX 4090 cards or a single L40S card. A high-end deployment could involve an A100, H100, or H200 GPU in a dedicated workstation or server [33](#) . The purchase cost for an H100 GPU alone can exceed \$15,000 USD. The Year 1 TCO would include the hardware purchase, server chassis, power supplies, cooling, and setup. Power consumption for a single H100 under load can be over 450 watts, requiring a robust electrical circuit and adequate cooling infrastructure, adding several thousand dollars to the initial setup cost [33](#) . Maintenance contracts and eventual hardware refresh cycles (every 3-5 years) must also be factored in. For a company with only 10-15 users, the high CapEx and limited scalability make this the least practical option for achieving the best ROI. The Year 3 TCO would primarily consist of maintenance and depreciation, but the initial investment remains substantial.

The following tables project the 2026 costs for key deployment scenarios. These figures are estimates based on current trends and publicly available pricing, adjusted for the Indian market.

**Table 1: Annual Cost Comparison (INR)**

| Deployment Option     | Model Example | Monthly Cost (INR)  | Annual Cost (INR)     | 3-Year TCO (INR)          |
|-----------------------|---------------|---------------------|-----------------------|---------------------------|
| Big-Tech API          | Claude Opus   | 15,000 - 40,000     | 1,80,000 - 4,80,000   | 5,40,000 - 14,40,000      |
| Cloud Self-Hosted     | Qwen 3 72B    | 1,50,000 - 5,50,000 | 18,00,000 - 66,00,000 | 54,00,000 - 198,00,000    |
| On-Premise Small      | Llama 4 Scout | 1,00,000 - 1,50,000 | 12,00,000 - 18,00,000 | 36,00,000 - 54,00,000     |
| On-Premise Enterprise | Qwen 3 72B    | 5,00,000 - 8,00,000 | 60,00,000 - 96,00,000 | 1,80,00,000 - 2,88,00,000 |

Table 2: Hardware Sizing and TCO for On-Premise Deployment

| Hardware              | GPU Type    | VRAM   | Purchase Cost (USD) | Estimated Server Cost (USD) | Year 1 TCO (USD) | Year 3 TCO (USD) |
|-----------------------|-------------|--------|---------------------|-----------------------------|------------------|------------------|
| Small Deployment      | RTX 4090 x2 | 48 GB  | 2,500               | 1,000                       | 4,500            | 4,000            |
| Small Deployment      | L40S        | 48 GB  | 12,000              | 2,500                       | 16,500           | 14,500           |
| Enterprise Deployment | A100 80GB   | 80 GB  | 15,000              | 3,000                       | 20,000           | 17,000           |
| Enterprise Deployment | H100 80GB   | 80 GB  | 18,000              | 3,500                       | 23,500           | 20,500           |
| Enterprise Deployment | H200        | 128 GB | 25,000              | 4,000                       | 31,500           | 28,000           |

*Note: Year 1 TCO assumes 500 USD for power/cooling/maintenance. Year 3 TCO assumes 200 USD for maintenance and depreciation. Costs are illustrative and do not include software licensing or personnel.*

This analysis clearly indicates that while APIs are simple, their long-term cost is prohibitive. On-premise offers the lowest marginal cost post-purchase but requires a very large upfront investment that is difficult to justify for this scale of operation. Cloud self-hosting emerges as the most balanced and economically rational path, offering a predictable, amortized cost structure over three years while providing the necessary control over data and security.

## The Impact of Indian Data Governance on System Design and Security

While the enterprise operates without a strict, government-mandated data residency requirement, the evolving regulatory landscape in India, particularly the Digital Personal

Data Protection (DPDP) Act, 2023, imposes a strong ethical and practical imperative to design systems with privacy and security at their core [2](#) [84](#). Building a system that aligns with the principles of the DPDP Act is not merely a compliance exercise; it is a strategic move towards creating a trustworthy, resilient, and future-proof platform. The architecture must therefore be viewed through the lens of "Privacy by Design" and "Security by Design," principles that are central to India's emerging AI governance framework [39](#) [64](#).

The DPDP Act establishes several foundational principles that directly influence system architecture. First is the principle of consent-based processing, which mandates that personal data can only be processed for a lawful, specific, and explicit purpose with the individual's consent [85](#) [88](#). Although the documents are internal company information, the underlying concept of purpose limitation is equally applicable. The ingestion pipeline must be designed to handle only the data necessary for the specified Q&A and analysis tasks, avoiding unnecessary collection and replication of information [39](#). Second, the Act enshrines the principle of data minimization, requiring that data collection be limited to what is adequate, relevant, and necessary for the stated purpose [88](#). This reinforces the need for a clean, efficient ingestion process that extracts only essential information from the PDFs. Third, the Act mandates robust security safeguards to protect personal data from unauthorized access, disclosure, and other harmful acts [88](#). This necessitates the implementation of end-to-end encryption for all data, whether at rest in the vector database and storage, or in transit between components of the system [39](#).

Furthermore, the DPDP Act places a heavy emphasis on accountability and transparency. Organizations must maintain a 'Record of Processing Activities' (RoPA) and conduct Data Protection Impact Assessments (DPIAs) for high-risk processing activities [85](#) [88](#). For an AI system, this translates into a non-negotiable requirement for immutable audit logging. Every action related to data processing—from document ingestion and chunking to query execution and result generation—must be logged in a tamper-evident manner. This creates a clear data lineage and provides the forensic trail necessary to demonstrate compliance to regulators and internal auditors [39](#). The Reserve Bank of India's (RBI) FREE-AI framework similarly emphasizes the need for transparency, oversight, and auditability across the entire AI lifecycle [2](#). An architecture that produces citable answers, as enabled by a well-designed RAG pipeline, directly supports this need for traceability and defensible decision logic [59](#).

India's AI governance guidelines, issued by the Ministry of Electronics and Information Technology (MeitY), complement the DPDP Act by outlining a broader framework for responsible AI [1](#) [109](#). These guidelines stress seven core principles, including safety,

security, sustainability, transparency, and accountability <sup>2</sup>. They advocate for a multi-tiered governance structure involving board-level oversight, dedicated ethics committees, and continuous assurance through testing <sup>64</sup>. From an architectural standpoint, this means the system should not be a "black box." It should incorporate mechanisms for monitoring model performance, detecting biases, and ensuring adversarial robustness <sup>64</sup>. The proposed mandatory licensing and royalty regime for text and data mining could also impact how models are trained, favoring approaches that rely less on scraped, unlicensed data <sup>2</sup>. Finally, while there are no immediate data localization laws, the DPDP Act does impose restrictions on cross-border data transfers, requiring Standard Contractual Clauses (SCCs) and annual reporting <sup>88</sup>. This makes self-hosted or private cloud deployments, which keep data within a single jurisdiction, a strategically sound choice to simplify compliance and mitigate transfer-related risks.

In summary, designing the PDF question-answering platform with Indian data governance in mind is not an optional extra; it is a fundamental requirement for building an enterprise-grade solution. By proactively integrating Privacy by Design, implementing end-to-end encryption, and establishing comprehensive audit trails, the enterprise can build a system that is not only functional and cost-effective but also secure, transparent, and compliant with the spirit of India's emerging data protection laws.

## Weighted Evaluation and Identification of Pareto-Optimal Solutions

To objectively compare the myriad of deployment options and identify the optimal architecture, a weighted scoring model is applied to every viable solution across eight key criteria. The weights reflect the user's prioritized goals: Accuracy (20%), Legal Analysis Quality (15%), Compliance Analysis Quality (15%), Hallucination Resistance (15%), Cost Efficiency (15%), Latency (5%), Security (5%), and Scalability (5%). Scores are assigned on a 1-10 scale, with 10 being the highest possible score for a given criterion. This rigorous evaluation allows for the identification of the Pareto Frontier—the set of solutions that are not dominated by any other option on all criteria—and ultimately leads to a single, justified final recommendation.

The evaluation considers three primary architectural paths: Big-Tech API, Cloud Self-Hosted, and On-Premise, each paired with a representative model and deployment strategy.

- 1. **Big-Tech API (e.g., Claude Opus via API):** This option scores highly on Operational Simplicity and Scalability due to its managed nature. However, it scores poorly on Cost Efficiency and Security because of the recurring marginal cost per query and the inherent risks of sending data to a third-party service. Its scores for Accuracy and Hallucination Resistance are good (8-9/10) due to the model's quality, but it cannot be optimized for legal/compliance niches. The lack of a self-hosted environment makes achieving perfect scores on Compliance Analysis Quality and Security difficult.
- 2. **Cloud Self-Hosted (e.g., Qwen 3 72B on AWS):** This path represents the most promising balance. It scores very highly on Accuracy, Legal Analysis Quality, and Hallucination Resistance (9/10) by leveraging a powerful open-source model that can be fine-tuned and deployed privately. Security is maximized (9/10) by keeping all data within a private cloud VPC. Cost Efficiency is good (7/10) after the initial setup, representing a favorable TCO over three years. Latency is achievable (<5s) with proper GPU sizing. This option forms the core of the Pareto Frontier.
- 3. **On-Premise (e.g., Qwen 3 72B on an H100 server):** This option achieves the highest possible scores for Security (10/10) and Cost Efficiency post-purchase (though initial CapEx is high). Its performance metrics (Accuracy, Hallucination Resistance) are comparable to the cloud self-hosted option (9/10). However, its scores for Scalability and Operational Simplicity are low (4/5) due to the physical limitations of the hardware and the need for in-house expertise to manage it. It is included in the Pareto Frontier for maximum security and lowest long-term marginal cost.

The following table summarizes the weighted scores for these key options. The scores are indicative and based on the synthesis of all provided sources.

| Solution                 | Accuracy | Legal/<br>Compliance<br>Quality | Hallucination<br>Resistance | Cost<br>Efficiency | Latency | Security | Scalability | Operational<br>Complexity | Weighted<br>Score |
|--------------------------|----------|---------------------------------|-----------------------------|--------------------|---------|----------|-------------|---------------------------|-------------------|
| Big-Tech<br>API          | 9        | 7                               | 8                           | 4                  | 8       | 6        | 9           | 3                         | 6.85              |
| Cloud<br>Self-<br>Hosted | 9        | 9                               | 9                           | 7                  | 8       | 9        | 7           | 6                         | 8.15              |
| On-<br>Premise           | 9        | 9                               | 9                           | 8                  | 7       | 10       | 4           | 5                         | 8.05              |

Based on this weighted evaluation, the Pareto-optimal solutions are the **Cloud Self-Hosted** and **On-Premise** architectures. Neither is universally superior; the Cloud Self-Hosted option offers better scalability and operational convenience, while the On-Premise option provides absolute data sovereignty and the lowest long-term marginal cost. The "Best Low-Cost Option" is the On-Premise architecture, assuming the initial capital investment is feasible. The "Best Accuracy Option" and "Best Legal Analysis Option" are tied between the Cloud Self-Hosted and On-Premise options, as both leverage powerful, privately deployed models. The "Best Security Option" is the On-Premise solution. The "Best Cloud Option" is the Cloud Self-Hosted architecture. The "Best ROI Option" is also strongly argued for the Cloud Self-Hosted model due to its balanced TCO and performance profile over a three-year period.

The final decision rests on the enterprise's appetite for capital expenditure versus operational expenditure and its tolerance for managing its own infrastructure. Given the company's small size and the desire for a balanced solution, the Cloud Self-Hosted architecture emerges as the most compelling choice, sitting squarely on the Pareto frontier and offering the best overall value proposition.

## Final Recommendation: The Balanced, Cost-Effective Hybrid Architecture

After a comprehensive analysis of the competitive landscape, architectural blueprints, deployment modalities, and regulatory considerations, the single most optimal solution for this India-based enterprise is a **Balanced, Cost-Effective Hybrid Architecture centered on a self-hosted, fine-tuned open-source language model**. This recommendation is grounded in the user's weighted evaluation framework, which prioritizes quality and reliability above all else, while still demanding cost-efficiency and

a pragmatic approach to operational complexity. The proposed architecture sits firmly on the Pareto frontier, offering the best possible trade-off between performance, security, and long-term cost.

The core of this recommended architecture is the deployment of a highly capable open-source model, such as **Qwen 3 72B** or **Llama 4 Maverick**, on a major cloud provider's GPU instances (e.g., AWS, GCP, or Azure) or a specialized, cost-effective provider like RunPod or Lambda Labs [54](#) [91](#) . This choice directly addresses the highest-weighted criteria in the evaluation. Open-source models now offer performance that rivals or exceeds proprietary equivalents in many domains, and their ability to be fine-tuned on the enterprise's specific corpus of legal and compliance documents will yield superior results for the "Legal Analysis Quality" and "Compliance Analysis Quality" metrics [71](#) [72](#) . This private deployment is critical for maximizing the "Security" score, as it ensures that sensitive internal documents never leave the organization's controlled cloud environment, thereby mitigating the risks associated with third-party APIs and aligning with the principles of the DPDP Act [39](#) [64](#) .

This architecture is projected to deliver the required sub-10-second latency for queries from 10-15 concurrent users by appropriately sizing the GPU resources, likely involving a small cluster of high-performance GPUs like A100s or H100s [33](#) . The Total Cost of Ownership (TCO) over a three-year period is estimated to be significantly lower than relying solely on Big-Tech APIs, making it the most cost-effective solution that still meets the stringent quality requirements [22](#) . This positions it as the "Best ROI Option" and "Best Future-Proof Option," as the investment in the model and infrastructure can be leveraged and expanded upon as the business grows. The system would employ a sophisticated RAG pipeline featuring advanced chunking, a hybrid retrieval strategy with re-ranking, and a citation-enforced generation process to maximize accuracy and minimize hallucination, directly targeting the 15% weight given to "Hallucination Resistance" [59](#) [93](#) .

While the pure On-Premise option offers marginally better security and the lowest long-term marginal cost, its prohibitive upfront capital expenditure and operational rigidity make it unsuitable for a small enterprise of this scale. The Big-Tech API option, while simple, fails to meet the requirements for cost-efficiency and deep security control at the projected query volume. Therefore, the self-hosted cloud solution provides the ideal equilibrium. It delivers the high-fidelity, trustworthy, and auditable performance demanded by the enterprise's use cases, all while maintaining a predictable and manageable financial profile. This architecture is not just a technical solution; it is a strategic foundation for responsible and effective enterprise AI in the Indian context.

---

## Reference

1. MeitY's AI Guidelines & The Evolving Copyright Landscape - Lexology <https://www.lexology.com/library/detail.aspx?g=ffc0c58c-3727-4472-9914-5fa6a33fffd>
2. AI Regulation in India: What to Expect in 2026 - LinkedIn <https://www.linkedin.com/pulse/ai-regulation-india-what-expect-2026-deepthi-rajeev-im6wc>
3. Gemini 3 Raises the Bar on Quality, But Not on Speed <https://www.legalontech.com/post/gemini-3-raises-the-bar-on-quality-but-not-on-speed>
4. Evaluating What Large Language Models (Actually) Know - arXiv <https://arxiv.org/html/2605.26937v1>
5. [PDF] Retrieval-Augmented Generation and Hallucination Span Detection <https://aclanthology.org/2025.semeval-1.151.pdf>
6. Real-Time Evaluation Models for RAG: Who Detects Hallucinations ... [https://www.researchgate.net/publication/390247766\\_Real-Time\\_Evaluation\\_Models\\_for\\_RAG\\_Who\\_Detects\\_Hallucinations\\_Best](https://www.researchgate.net/publication/390247766_Real-Time_Evaluation_Models_for_RAG_Who_Detects_Hallucinations_Best)
7. Real-Time Evaluation Models for RAG: Who Detects Hallucinations ... <https://arxiv.org/abs/2503.21157>
8. DefAn: Definitive Answer Dataset for LLM Hallucination Evaluation <https://www.mdpi.com/2078-2489/16/11/937>
9. [PDF] HaluEval: A Large-Scale Hallucination Evaluation Benchmark ... <https://www.semanticscholar.org/paper/HaluEval%3A-A-Large-Scale-Hallucination-Evaluation-Li-Cheng/e0384ba36555232c587d4a80d527895a095a9001>
10. Detecting RAG Hallucinations: A 2025 Enterprise Guide - LinkedIn <https://www.linkedin.com/pulse/detecting-rag-hallucinations-2025-enterprise-guide-lindegaard-uvdmf>
11. [PDF] ODE: Open-Set Evaluation of Hallucinations in Multimodal Large ... [https://openaccess.thecvf.com/content/CVPR2025/papers/Tu\\_ODE\\_Open-Set\\_Evaluation\\_of\\_Hallucinations\\_in\\_Multimodal\\_Large\\_Language\\_Models\\_CVPR\\_2025\\_paper.pdf](https://openaccess.thecvf.com/content/CVPR2025/papers/Tu_ODE_Open-Set_Evaluation_of_Hallucinations_in_Multimodal_Large_Language_Models_CVPR_2025_paper.pdf)
12. Sovereign cloud implementation: Technical architectures for data ... [https://www.researchgate.net/publication/390466414\\_Sovereign\\_cloud\\_implementation\\_Technical\\_architectures\\_for\\_data\\_residency\\_and\\_regulatory\\_compliance](https://www.researchgate.net/publication/390466414_Sovereign_cloud_implementation_Technical_architectures_for_data_residency_and_regulatory_compliance)
13. Amazon API Gateway Pricing - AWS <https://aws.amazon.com/api-gateway/pricing/>

14. Amazon Connect Customer Pricing - AWS <https://aws.amazon.com/products/connect/customer/pricing/>
15. AWS Product and Service Pricing - Amazon.com <https://aws.amazon.com/pricing/>
16. Amazon S3 pricing - AWS <https://aws.amazon.com/s3/pricing/>
17. AI Coding & Dev Tools - Amazon Q - AWS <https://aws.amazon.com/q/>
18. Amazon Textract pricing - AWS <https://aws.amazon.com/textract/pricing/>
19. Cloud Computing Services - Amazon Web Services (AWS) <https://aws.amazon.com/>
20. Calling AWS services and prices using the AWS Price List <https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/price-changes.html>
21. Amazon API Gateway | API Management | Amazon Web Services <https://aws.amazon.com/api-gateway/>
22. Amazon SageMaker AI Pricing <https://aws.amazon.com/sagemaker/ai/pricing/>
23. A Comprehensive Review of the Architecture and Trust Frameworks ... <https://arxiv.org/html/2601.05264>
24. Benchmarking Complex Multimodal Document Processing Pipelines <https://arxiv.org/html/2604.26382v1>
25. Large Language Models, Architectures, Protocols, and Applications <https://arxiv.org/html/2306.02781v4>
26. Federated Retrieval-Augmented Generation: A Systematic Mapping ... <https://arxiv.org/html/2505.18906v1>
27. Artificial Intelligence - arXiv <https://arxiv.org/list/cs.AI/new>
28. 1. Introduction - arXiv <https://arxiv.org/html/2603.04740v1>
29. From Logic Monopoly to Social Contract: Separation of Power and ... <https://arxiv.org/html/2603.25100v1>
30. AI Agents Under EU Law A Compliance Architecture for AI Providers <https://arxiv.org/html/2604.04604v1>
31. Governance by Design: Architecting Agentic AI for Organizational ... <https://arxiv.org/html/2605.20210v1>
32. Taxonomy and Consistency Analysis of Safety Benchmarks for AI ... <https://arxiv.org/html/2605.16282v1>
33. NVIDIA A100 Guide: How to Choose the Right GPU for AI Workloads <https://electronics.alibaba.com/buyingguides/nvidia-a100-guide-buy-vs-rent,-40gb-vs-80gb>
34. A list of large language models (LLMs) - IBM <https://www.ibm.com/think/topics/large-language-models-list>

35. AI Governance Playbook: Enterprise Framework <https://petronellatech.com/blog/ai-governance-playbook-model-risk-compliance-and-scalable-automation/>
36. AI Agents vs. Agentic AI: A Conceptual taxonomy, applications and ... <https://www.sciencedirect.com/science/article/pii/S1566253525006712>
37. KYA: A Framework-Agnostic Trust Layer for Autonomous Systems ... <https://arxiv.org/html/2605.25376v1>
38. Cloud-Driven Security Compliance: Architecting GDPR & CCPA ... [https://www.researchgate.net/publication/395477294\\_Cloud-Driven\\_Security\\_Compliance\\_Architecting\\_GDPR\\_CCPA\\_Solutions\\_For\\_Large-Scale\\_Digital\\_Platforms](https://www.researchgate.net/publication/395477294_Cloud-Driven_Security_Compliance_Architecting_GDPR_CCPA_Solutions_For_Large-Scale_Digital_Platforms)
39. How to Build a DPDP-Compliant AI Data Architecture <https://jisasoftech.com/how-to-build-a-dpdp-compliant-ai-data-architecture/>
40. Artificial intelligence readiness assessment report: India <https://unesdoc.unesco.org/ark:/48223/pf0000397347>
41. [PDF] AI for ALL: Catalysing Jobs, Growth, and Opportunity <https://web-assets.bcg.com/fa/7d/2034b2d04339aa9d94d70a9be759/ai-for-all.pdf>
42. Trending Applications of Large Language Models - IEEE Xplore <https://ieeexplore.ieee.org/iel8/9078688/9184921/11199892.pdf>
43. OpenAI's GPT-5.2 Falls Short, Focusing on Benchmarks ... - LinkedIn [https://www.linkedin.com/posts/andrewamann\\_openai-is-04-on-agi-attempts-this-year-activity-7408915950772584450-43wy](https://www.linkedin.com/posts/andrewamann_openai-is-04-on-agi-attempts-this-year-activity-7408915950772584450-43wy)
44. [PDF] SafeLawBench: Towards Safe Alignment of Large Language Models <https://aclanthology.org/2025.findings-acl.721.pdf>
45. [PDF] AgriEval: A Comprehensive Chinese Agricultural Benchmark for ... <https://ojs.aaai.org/index.php/AAAI/article/view/40716/44677>
46. LEXam: Benchmarking Legal Reasoning on 340 Law Exams - arXiv <https://arxiv.org/html/2505.12864v1>
47. On Verifiable Legal Reasoning: A Multi-Agent Framework with ... <https://dl.acm.org/doi/pdf/10.1145/3746252.3761057>
48. Sharad Bagri - Software Engineer at Google - LinkedIn <https://www.linkedin.com/in/sharadbags>
49. Managing your payments in India - AWS Billing - AWS Documentation <https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/edit-aispl-payment-method.html>
50. [PDF] AWS Marketplace - Buyer Guide <https://docs.aws.amazon.com/pdfs/marketplace/latest/buyerguide/aws-marketplace-bug.pdf>

51. Sarvam launches on AWS Marketplace in India - LinkedIn [https://www.linkedin.com/posts/sarvam-ai\\_were-thrilled-to-share-that-sarvam-is-now-activity-7392534537110511616-zZQW](https://www.linkedin.com/posts/sarvam-ai_were-thrilled-to-share-that-sarvam-is-now-activity-7392534537110511616-zZQW)
52. Estimating Black-Box LLM Parameter Counts via Factual Capacity <https://arxiv.org/html/2604.24827v1>
53. RunPod vs Lambda Labs vs CoreWeave: GPU Cloud Pricing <https://www.buildmvpfast.com/blog/gpu-cloud-cost-comparison-runpod-lambda-labs-coreweave-2026>
54. A Comparative Study of Cloud GPU Offerings for AI/ML Engineers [https://www.researchgate.net/publication/403638287\\_A\\_Comparative\\_Study\\_of\\_Cloud\\_GPU\\_Offerings\\_for\\_AIML\\_Engineers](https://www.researchgate.net/publication/403638287_A_Comparative_Study_of_Cloud_GPU_Offerings_for_AIML_Engineers)
55. HalluHard: A Hard Multi-Turn Hallucination Benchmark - arXiv <https://arxiv.org/html/2602.01031v1>
56. DescrybeLM Outperforms ChatGPT, Claude on Bar Exam - LinkedIn [https://www.linkedin.com/posts/robertambrogi\\_ai-legal-research-startup-descrybe-launches-activity-7435339747633377280-4O6G](https://www.linkedin.com/posts/robertambrogi_ai-legal-research-startup-descrybe-launches-activity-7435339747633377280-4O6G)
57. Hallucination – Free? Assessing the Reliability of Leading AI Legal ... [https://www.researchgate.net/publication/391086271\\_Hallucination-Free\\_Assessing\\_the\\_Reliability\\_of\\_Leading\\_AI\\_Legal\\_Research\\_Tools](https://www.researchgate.net/publication/391086271_Hallucination-Free_Assessing_the_Reliability_of_Leading_AI_Legal_Research_Tools)
58. Best Vector Databases for RAG: Top Picks for 2026 - TiDB <https://www.pingcap.com/compare/best-vector-database/>
59. Citation-Enforced RAG for Fiscal Document Intelligence - arXiv <https://arxiv.org/html/2603.14170v1>
60. RAG Enterprise Implementation, Governance and Evaluation <https://www.linkedin.com/pulse/retrieval-augmented-generation-rag-part-2-hany-tadros-ajcee>
61. The State of LLMs in 2025: GPT-5, Claude, Gemini, Llama 4 ... <https://www.linkedin.com/pulse/state-llms-2025-gpt-5-claude-gemini-llama-4-deepseek-complete-keyan-yidbc>
62. [PDF] India AI Governance Guidelines - AWS <https://indiaai.s3.ap-south-1.amazonaws.com/docs/guidelines-governance.pdf>
63. AI Governance Guidelines for Businesses | PDF - Scribd <https://www.scribd.com/document/836118213/AI-Guidelines-Whitepaper>
64. AI Governance Framework for India 2025-26 - LinkedIn <https://www.linkedin.com/pulse/ai-governance-framework-india-2025-26-ninad-chavan-csj1f>
65. Abstract - arXiv <https://arxiv.org/html/2509.10707v1>

66. [PDF] AI Privacy Risks & Mitigations – Large Language Models (LLMs) <https://www.edpb.europa.eu/system/files/2025-04/ai-privacy-risks-and-mitigations-in-llms.pdf>
67. A Systematic Review of Key Retrieval-Augmented Generation (RAG ... <https://arxiv.org/html/2507.18910v1>
68. Large Language Models Hallucination: A Comprehensive Survey <https://arxiv.org/html/2510.06265v3>
69. [PDF] Large Language Models Hallucination: A Comprehensive Survey <https://arxiv.org/pdf/2510.06265>
70. Part 3: India AI Governance Guidelines || Implementation Roadmap <https://www.linkedin.com/pulse/part-3-india-ai-governance-guidelines-implementation-roadmap-shah-xejuf>
71. AI Built for Law Outperforms ChatGPT, Claude, and Gemini on Legal ... <https://www.stargazette.com/press-release/story/41920/ai-built-for-law-outperforms-chatgpt-claude-and-gemini-on-legal-reasoning-benchmark/>
72. [PDF] LawBench: Benchmarking Legal Knowledge of Large Language ... <https://arxiv.org/pdf/2309.16289>
73. Analysis of solutions for a blockchain compliance with GDPR - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC9440070/>
74. 2026年全网最全大模型API横评：Claude / GPT / Gemini 等8 大厂商 ... <https://segmentfault.com/a/1190000047676047>
75. [PDF] A Comparison of DeepSeek, Qwen, ChatGPT, Gemini, Llama ... - arXiv <https://arxiv.org/pdf/2503.04765>
76. [PDF] HalluHard: A Hard Multi-Turn Hallucination Benchmark - arXiv <https://arxiv.org/pdf/2602.01031?>
77. [PDF] Evaluating Evaluation Metrics – The Mirage of Hallucination Detection <https://aclanthology.org/2025.findings-emnlp.1035.pdf>
78. HaluEval: A Large-Scale Hallucination Evaluation Benchmark for ... [https://www.researchgate.net/publication/376394138\\_HaluEval\\_A\\_Large-Scale\\_Hallucination\\_Evaluation\\_Benchmark\\_for\\_Large\\_Language\\_Models](https://www.researchgate.net/publication/376394138_HaluEval_A_Large-Scale_Hallucination_Evaluation_Benchmark_for_Large_Language_Models)
79. HaluEval - OpenDataLab <https://opendatalab.com/OpenDataLab/HaluEval/download>
80. [PDF] TriDF: Evaluating Perception, Detection, and Hallucination for ... [https://openaccess.thecvf.com/content/CVPR2026/papers/Jiang-Lin\\_TriDF\\_Evaluating\\_Perception\\_Detection\\_and\\_Hallucination\\_for\\_Interpretable\\_DeepFake\\_Detection\\_CVPR\\_2026\\_paper.pdf](https://openaccess.thecvf.com/content/CVPR2026/papers/Jiang-Lin_TriDF_Evaluating_Perception_Detection_and_Hallucination_for_Interpretable_DeepFake_Detection_CVPR_2026_paper.pdf)
81. 2025 LLM Knowledge & Reasoning Benchmarks | PDF - Scribd <https://www.scribd.com/document/966797651/LLM-Benchmark>

82. Automated Policy Evaluation and Legal Document Analysis [https://www.researchgate.net/publication/393978992\\_Enhancing\\_GDPR\\_Compliance\\_Through\\_NLP\\_Automated\\_Policy\\_Evaluation\\_and\\_Legal\\_Document\\_Analysis](https://www.researchgate.net/publication/393978992_Enhancing_GDPR_Compliance_Through_NLP_Automated_Policy_Evaluation_and_Legal_Document_Analysis)
83. An Empirical Analysis of Static Analysis Methods for Detection and ... <https://arxiv.org/html/2604.07755v2>
84. AI transforms compliance with DPDP rules in India - LinkedIn [https://www.linkedin.com/posts/vasanth-kumar-9b1360257\\_compliance-ai-audit-activity-7374504360371785728-RbfJ](https://www.linkedin.com/posts/vasanth-kumar-9b1360257_compliance-ai-audit-activity-7374504360371785728-RbfJ)
85. India's DPDP Act 2023: Security Implications for Architects - LinkedIn [https://www.linkedin.com/posts/ajay-singh-baghel-6175a086\\_cybersecurity-dpdp-dataprivacy-activity-7438879735218130944-euek](https://www.linkedin.com/posts/ajay-singh-baghel-6175a086_cybersecurity-dpdp-dataprivacy-activity-7438879735218130944-euek)
86. (PDF) Retrieval-Augmented Generation (RAG) and Large Language ... [https://www.researchgate.net/publication/399264692\\_Retrieval-Augmented\\_Generation\\_RAG\\_and\\_Large\\_Language\\_Models\\_LLMs\\_for\\_Enterprise\\_Knowledge\\_Management\\_and\\_Document\\_Automation\\_A\\_Systematic\\_Literature\\_Review](https://www.researchgate.net/publication/399264692_Retrieval-Augmented_Generation_RAG_and_Large_Language_Models_LLMs_for_Enterprise_Knowledge_Management_and_Document_Automation_A_Systematic_Literature_Review)
87. Gaps or Hallucinations? Scrutinizing Machine-Generated Legal ... <https://aclanthology.org/2024.nllp-1.24/>
88. Detailed Comparative Structure Between DPDP Act (2023) and ... <https://www.linkedin.com/pulse/detailed-comparative-structure-between-dpdp-act-2023-gdpr-policy-t5ljc>
89. Compliance Obligations under the Digital Personal Data Protection ... [https://www.researchgate.net/publication/401631009\\_Compliance\\_Obligations\\_under\\_the\\_Digital\\_Personal\\_Data\\_Protection\\_Act\\_2023\\_A\\_Critical\\_and\\_Comparative\\_Analysis\\_with\\_the\\_GDPR](https://www.researchgate.net/publication/401631009_Compliance_Obligations_under_the_Digital_Personal_Data_Protection_Act_2023_A_Critical_and_Comparative_Analysis_with_the_GDPR)
90. [PDF] Enhancing Legal Compliance and Regulation Analysis with Large ... <https://arxiv.org/pdf/2404.17522>
91. Lambda vs Thunder Compute (2026): Pricing, Billing, and Experience <https://www.thundercompute.com/blog/lambda-labs-vs-thunder-compute>
92. An Auditable LLM-RAG Architecture for Financial Document ... - MDPI <https://www.mdpi.com/1999-5903/18/6/284>
93. [2603.14170] Citation-Enforced RAG for Fiscal Document Intelligence <https://arxiv.org/abs/2603.14170>
94. OmniCompliance-100K: A Multi-Domain, Rule-Grounded, Real ... <https://arxiv.org/html/2603.13933v2>
95. Uncovering Vulnerabilities of LLM-Assisted Cyber Threat Intelligence <https://arxiv.org/html/2509.23573v5>

96. VeriGuard: Enhancing LLM Agent Safety via Verified Code Generation <https://arxiv.org/html/2510.05156v1>
97. [PDF] COMPL-AI Framework: A Technical Interpretation and LLM ... - arXiv <https://arxiv.org/pdf/2410.07959>
98. LLMOrbit: A Circular Taxonomy of Large Language Models - arXiv <https://arxiv.org/html/2601.14053v1>
99. A Comprehensive Survey on Safety Evaluation of LLMs - arXiv <https://arxiv.org/html/2506.11094v1>
100. Privasis: Synthesizing the Largest “Public” Private Dataset ... - arXiv <https://arxiv.org/html/2602.03183v1>
101. Training AI Co-Scientists Using Rubric Rewards - arXiv <https://arxiv.org/html/2512.23707v1>
102. FlowSteer: Interactive Agentic Workflow Orchestration via End-to ... <https://arxiv.org/html/2602.01664v1>
103. [PDF] HalluHard: A Hard Multi-Turn Hallucination Benchmark - arXiv <https://arxiv.org/pdf/2602.01031>
104. [PDF] Exploring the Generalizability of Factual Hallucination Mitigation via ... <https://aclanthology.org/2025.findings-emnlp.211.pdf>
105. (PDF) The Most Advanced AI Models of 2025 -Comparative Analysis ... [https://www.researchgate.net/publication/392160200\\_The\\_Most\\_Advanced\\_AI\\_Models\\_of\\_2025\\_-\\_Comparative\\_Analysis\\_of\\_Gemini\\_25\\_Claude\\_4\\_LLaMA\\_4\\_GPT-45\\_DeepSeek\\_V31\\_and\\_Other\\_Leading\\_Models](https://www.researchgate.net/publication/392160200_The_Most_Advanced_AI_Models_of_2025_-_Comparative_Analysis_of_Gemini_25_Claude_4_LLaMA_4_GPT-45_DeepSeek_V31_and_Other_Leading_Models)
106. Safeguarding large language models: a survey - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC12532640/>
107. On large language models safety, security, and privacy: A survey <https://www.sciencedirect.com/science/article/pii/S1674862X25000023>
108. SWE-bench Leaderboards <https://www.swebench.com/>
109. Summary of recently published "India AI Governance Guidelines <https://www.linkedin.com/pulse/summary-recently-published-india-ai-governance-safe-xfkqf>
110. A New In-Depth Report of AI Large Language Models: Hallucination ... <https://www.hku.hk/press/press-releases/detail/28587.html>
111. [PDF] HalluLens: LLM Hallucination Benchmark - ACL Anthology <https://aclanthology.org/2025.acl-long.1176.pdf>
112. Zero-knowledge LLM hallucination detection and mitigation through ... <https://arxiv.org/html/2508.14314v1>

113. Survey and analysis of hallucinations in large language models - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC12518350/>
114. An Empirical Study on Factuality Hallucination in Large Language ... <https://aclanthology.org/2024.acl-long.586/>